

Implementing a Simple Q-Learning Model

Q-Learning is a **model-free reinforcement learning algorithm** used to learn the value of actions in a given environment. The agent updates a **Q-table** to estimate the expected cumulative reward for each state-action pair.

Key Components of Q-Learning

1. **Q-Table:** Stores values of state-action pairs.
 2. **Learning Rate (α):** Determines how much new information overrides old.
 3. **Discount Factor (γ):** Balances immediate and future rewards.
 4. **Exploration Rate (ϵ):** Probability of taking random actions to explore the environment.
-

Algorithm Steps

1. Initialize Q-table with zeros.
2. For each episode:
 - Observe initial state.
 - Select an action (ϵ -greedy: explore or exploit).
 - Execute action, observe reward and next state.
 - Update Q-value using the formula:

$$Q(s,a) = Q(s,a) + \alpha [r + \gamma \max_{a'} Q(s',a') - Q(s,a)]$$
$$Q(s,a) = Q(s,a) + \alpha [r + \gamma \max_{a'} Q(s',a') - Q(s,a)]$$

- Repeat until reaching terminal state.
-

Python Example Using OpenAI Gym (FrozenLake Environment):

```
import gym

import numpy as np

# Initialize environment

env = gym.make('FrozenLake-v1', is_slippery=False)
```

```

n_states = env.observation_space.n

n_actions = env.action_space.n

q_table = np.zeros((n_states, n_actions))


# Hyperparameters

alpha = 0.8    # Learning rate

gamma = 0.95   # Discount factor

epsilon = 0.1  # Exploration rate

episodes = 1000


# Q-Learning algorithm

for episode in range(episodes):

    state = env.reset()

    done = False

    while not done:

        # Choose action ( $\epsilon$ -greedy)

        if np.random.rand() < epsilon:

            action = env.action_space.sample() # Explore

        else:

            action = np.argmax(q_table[state]) # Exploit


        # Take action

        next_state, reward, done, _ = env.step(action)


        # Update Q-value

        q_table[state, action] = q_table[state, action] + alpha * (

            reward + gamma * np.max(q_table[next_state]) - q_table[state, action]

        )

```

```
state = next_state

# Evaluate Q-table
print("Trained Q-Table:\n", q_table)

# Test the trained agent
state = env.reset()
done = False
while not done:
    action = np.argmax(q_table[state])
    state, reward, done, _ = env.step(action)
    env.render()
print("Reward:", reward)
```

Notes

- Q-Learning works best in **discrete environments** with a finite number of states and actions.
- For large or continuous state spaces, **Deep Q-Networks (DQN)** are used instead of Q-tables.
- Hyperparameters like learning rate, discount factor, and exploration rate significantly affect learning efficiency.

This simple Q-learning implementation provides a **foundation for reinforcement learning**, demonstrating how an agent can learn optimal actions by trial and error.